Lecture 2 - Tables / Insert / Select

Feb 20, 2020

SQL databases are composed of a set of tables.

A table is a set of columns, with names and types - like a struct or record in other languages.

There is a set of rows - that act kind of like an un-ordered array of columns.

SQL normally stores this as a row-order store.

Some SQL-ish databases like Cassandra store this as a colum-store.

```
$ psql
pg=# \i create-tables.sql
pg=# \q
$
```

the file

```
 1: -- \c l02
 2: create table vote_by_county (
 3:     id                serial primary key,
 4:     year             int default 2021,
 5:     state             text default '--',        -- irritatingly all upper case.
 6:     state_uc         text default '--',
 7:     state_po         varchar(2) default '--',     -- Incorrectly Named Column!
 8:     county_name         text default '--',        -- irritatingly all upper case.
 9:     county_name_uc    text default '--',
10:     county_fips         int default 0,
11:     office             text default 'unk',
12:     candidate         text default 'unk',
13:     candidate_uc     text default 'unk',
14:     party             text default 'unk',
15:     candidatevotes     int default 0,
16:     totalvotes         int default 0,
17:     version             int,
18:     vote_mode         text
19: );
```

Output

```
CREATE TABLE
```

Let's just insert a few rows to see how insert works:

```
1: -- \c l02
2:
3: insert into vote_by_county  ( year, state, county_name, version ) values
4:     ( 2022, 'Wyoming', 'Albeny',   1 );
5: insert into vote_by_county  ( year, state, county_name, version ) values
6:     ( 2022, 'Wyoming', 'Big Horn', 2 );
7: insert into vote_by_county  ( year, state, county_name, version ) values
8:     ( 2022, 'Wyoming', 'Carbon',   8 );
```

Output

```
INSERT 0 1
INSERT 0 1
INSERT 0 1
```

Now we can get the data back:

```
select year, state, county_name, version
        from vote_by_county ;
```

Data Types for tables

The data types used in our example:

Type	Description
serial	an integer that is generated (if null) by the database.
int	an integer (bigint, numeric etc)
text	a string from NULL, 0 length to `2**33` bytes in length.
varchar(maxlen)	A 'text' field with a limited length.
char(maxlen)	a right blank-padded fixed length string.
char varying(maxlen)	same as varchar()

`not null` disallow NULL values.

`default Value` if insert skips this value then default will be used. This is different than the value NULL.

`primary key` means that .

There are a bunch of other data types in PostgreSQL (MySQL/MariaDB has a bunch too, Oracle has a bunch, MS Sql Server too).

Name	Aliases	Description
`bigint`	`int8`	signed eight-byte integer
`bigserial`	`serial8`	autoincrementing eight-byte integer

Name	Aliases	Description
bit [(n)]		fixed-length bit string
bit varying [(n)]	varbit [(n)]	variable-length bit string
boolean	bool	logical Boolean (true/false)
box		rectangular box on a plane
bytea		binary data (“byte array”)
character [(n)]	char [(n)]	fixed-length character string
character varying [(n)]	varchar [(n)]	variable-length character string
cidr		IPv4 or IPv6 network address
circle		circle on a plane
date		calendar date (year, month, day)
double precision	float8	double precision floating-point number (8 bytes)
inet		IPv4 or IPv6 host address
integer	int, int4	signed four-byte integer
interval [fields] [(p)]		time span
json		textual JSON data
jsonb		binary JSON data, decomposed
line		infinite line on a plane
lseg		line segment on a plane
macaddr		MAC (Media Access Control) address
money		currency amount
numeric [(p, s)]	decimal [(p, s)]	exact numeric of selectable precision
path		geometric path on a plane
pg_lsn		PostgreSQL Log Sequence Number
point		geometric point on a plane
polygon		closed geometric path on a plane
real	float4	single precision floating-point number (4 bytes)
smallint	int2	signed two-byte integer
smallserial	serial2	autoincrementing two-byte integer
serial	serial4	autoincrementing four-byte integer
text		variable-length character string
time [(p)] [without time zone]		time of day (no time zone)
time [(p)] with time zone	timetz	time of day, including time zone

Name	Aliases	Description
`timestamp [ (p) ] [ without time zone ]`		date and time (no time zone)
`timestamp [ (p) ] with time zone`	`timestamptz`	date and time, including time zone
`tsquery`		text search query
`tsvector`		text search document
`txid_snapshot`		user-level transaction ID snapshot
`uuid`		universally unique identifier
`xml`		XML data