

Lecture 03 - Select

Jan 25

```
Select [ Projected Columns ]
From Table ...
    join to ...
Where ...
Group By Column List
Having [ where on grouped data ]
Ordr By [ Columns ]
```

Let's use our table from last time

```
1: -- \c l02
2: create table vote_by_county (
3:     id                serial primary key,
4:     year              int default 2021,
5:     state             text default '--',      -- irritatingly all upper case.
6:     state_uc          text default '--',
7:     state_po         varchar(2) default '--',  -- Incorrectly Named Column!
8:     county_name       text default '--',      -- irritatingly all upper case.
9:     county_name_uc   text default '--',
10:    county_fips        int default 0,
11:    office             text default 'unk',
12:    candidate          text default 'unk',
13:    candidate_uc      text default 'unk',
14:    party             text default 'unk',
15:    candidatevotes     int default 0,
16:    totalvotes        int default 0,
17:    version            int,
18:    vote_mode         text
19: );
```

Let's just insert a few rows to see how insert works:

```
1: -- \c l02
2:
3: insert into vote_by_county ( year, state, county_name, version ) values
4:     ( 2022, 'Wyoming', 'Albeny', 1 );
5: insert into vote_by_county ( year, state, county_name, version ) values
6:     ( 2022, 'Wyoming', 'Big Horn', 2 );
7: insert into vote_by_county ( year, state, county_name, version ) values
8:     ( 2022, 'Wyoming', 'Carbon', 8 );
```

and do some selects with the projected columns.

```
select id, year
from vote_by_county
;
```

we can rename a column

```
select id, year as "Year of Our Lord"
from vote_by_county
;
```

we can pick different columns

File: 01.sql

```
select id, year, state, county_name as "county"
from vote_by_county
;
```

Output:

id	year	state	county
16	2000	Alabama	Bibb
1992	2000	Georgia	Troup
4932	2000	Michigan	Ingham
5704	2000	Mississippi	Perry
12884	2004	Arkansas	Pope
...			
72592	2020	Wyoming	Sheridan
72593	2020	Wyoming	Sheridan
72594	2020	Wyoming	Sublette
72595	2020	Wyoming	Sublette
72596	2020	Wyoming	Sublette
72597	2020	Wyoming	Sublette
72598	2020	Wyoming	Sweetwater
72599	2020	Wyoming	Sweetwater
72600	2020	Wyoming	Sweetwater
72601	2020	Wyoming	Sweetwater
72602	2020	Wyoming	Teton
72603	2020	Wyoming	Teton
72604	2020	Wyoming	Teton
72605	2020	Wyoming	Teton
72606	2020	Wyoming	Uinta
72607	2020	Wyoming	Uinta
72608	2020	Wyoming	Uinta
72609	2020	Wyoming	Uinta
72610	2020	Wyoming	Washakie
72611	2020	Wyoming	Washakie
72612	2020	Wyoming	Washakie
72613	2020	Wyoming	Washakie
72614	2020	Wyoming	Weston
72615	2020	Wyoming	Weston
72616	2020	Wyoming	Weston
72617	2020	Wyoming	Weston

(72617 rows)

How about sorting the data

File: 02.sql

```
select id, year, state, county_name
  from vote_by_county
 order by county_name, state
;
```

Output:

id	year	state	county_name
28679	2008	South Carolina	Abbeville
38030	2012	South Carolina	Abbeville
38029	2012	South Carolina	Abbeville
38028	2012	South Carolina	Abbeville
28677	2008	South Carolina	Abbeville
9148	2000	South Carolina	Abbeville
9147	2000	South Carolina	Abbeville
9146	2000	South Carolina	Abbeville
9145	2000	South Carolina	Abbeville
19328	2004	South Carolina	Abbeville
19326	2004	South Carolina	Abbeville
19327	2004	South Carolina	Abbeville
28678	2008	South Carolina	Abbeville
...			
66634	2020	Texas	Zavala
66635	2020	Texas	Zavala
48759	2016	Texas	Zavala
66633	2020	Texas	Zavala
47714	2016	South Dakota	Ziebach
69518	2020	South Dakota	Ziebach
47713	2016	South Dakota	Ziebach
47712	2016	South Dakota	Ziebach
69516	2020	South Dakota	Ziebach
69517	2020	South Dakota	Ziebach
38362	2012	South Dakota	Ziebach
38363	2012	South Dakota	Ziebach
29012	2008	South Dakota	Ziebach
9589	2000	South Dakota	Ziebach
29011	2008	South Dakota	Ziebach
29010	2008	South Dakota	Ziebach
19661	2004	South Dakota	Ziebach
19660	2004	South Dakota	Ziebach
9590	2000	South Dakota	Ziebach
19659	2004	South Dakota	Ziebach
9591	2000	South Dakota	Ziebach
38361	2012	South Dakota	Ziebach
9592	2000	South Dakota	Ziebach

(72617 rows)

you can only sort by the columns that you have in the *projected columns*.

you can use the column position

File: 03.sql

```

select id, year, state, county_name
  from vote_by_county
 order by 4, 3
;

```

Output:

id	year	state	county_name
28679	2008	South Carolina	Abbeville
38030	2012	South Carolina	Abbeville
38029	2012	South Carolina	Abbeville
38028	2012	South Carolina	Abbeville
28677	2008	South Carolina	Abbeville
9148	2000	South Carolina	Abbeville
9147	2000	South Carolina	Abbeville
9146	2000	South Carolina	Abbeville
9145	2000	South Carolina	Abbeville
19328	2004	South Carolina	Abbeville
19326	2004	South Carolina	Abbeville
19327	2004	South Carolina	Abbeville
28678	2008	South Carolina	Abbeville
47379	2016	South Carolina	Abbeville
...			
38362	2012	South Dakota	Ziebach
38363	2012	South Dakota	Ziebach
29012	2008	South Dakota	Ziebach
9589	2000	South Dakota	Ziebach
29011	2008	South Dakota	Ziebach
29010	2008	South Dakota	Ziebach
19661	2004	South Dakota	Ziebach
19660	2004	South Dakota	Ziebach
9590	2000	South Dakota	Ziebach
19659	2004	South Dakota	Ziebach
9591	2000	South Dakota	Ziebach
38361	2012	South Dakota	Ziebach
9592	2000	South Dakota	Ziebach

(72617 rows)

you can ascending or descending sort

File: 04.sql

```

select id, year, state, county_name
  from vote_by_county
 order by 4 desc, 3 asc
;

```

Output:

id	year	state	county_name
9591	2000	South Dakota	Ziebach
19659	2004	South Dakota	Ziebach
19660	2004	South Dakota	Ziebach
19661	2004	South Dakota	Ziebach
38361	2012	South Dakota	Ziebach
38362	2012	South Dakota	Ziebach
38363	2012	South Dakota	Ziebach
29012	2008	South Dakota	Ziebach
29011	2008	South Dakota	Ziebach
9589	2000	South Dakota	Ziebach
29010	2008	South Dakota	Ziebach
9592	2000	South Dakota	Ziebach
9590	2000	South Dakota	Ziebach
47714	2016	South Dakota	Ziebach
47713	2016	South Dakota	Ziebach
47712	2016	South Dakota	Ziebach
69518	2020	South Dakota	Ziebach
69517	2020	South Dakota	Ziebach
69516	2020	South Dakota	Ziebach
48761	2016	Texas	Zavala
66637	2020	Texas	Zavala
66633	2020	Texas	Zavala
...			
9148	2000	South Carolina	Abbeville
38030	2012	South Carolina	Abbeville
38029	2012	South Carolina	Abbeville
38028	2012	South Carolina	Abbeville
28677	2008	South Carolina	Abbeville
28678	2008	South Carolina	Abbeville
28679	2008	South Carolina	Abbeville
19326	2004	South Carolina	Abbeville
19327	2004	South Carolina	Abbeville

(72617 rows)

You can apply functions and operators to the columns. In this case I will add 10 to the year and concatenate, || the state and county.

File: 05.sql

```
select id, year + 10 as "x", state||', '||county_name as "Location"
  from vote_by_county
 order by 3
;
```

Output:

id	x	Location
31169	2022	Alabama, Autauga
31168	2022	Alabama, Autauga
31167	2022	Alabama, Autauga
1	2010	Alabama, Autauga
21818	2018	Alabama, Autauga

```

21817 | 2018 | Alabama, Autauga
21816 | 2018 | Alabama, Autauga
12465 | 2014 | Alabama, Autauga
12466 | 2014 | Alabama, Autauga
12467 | 2014 | Alabama, Autauga
    4 | 2010 | Alabama, Autauga
    3 | 2010 | Alabama, Autauga
    2 | 2010 | Alabama, Autauga
40519 | 2026 | Alabama, Autauga
50527 | 2030 | Alabama, Autauga
40518 | 2026 | Alabama, Autauga
50525 | 2030 | Alabama, Autauga
...
49853 | 2026 | Wyoming, Uinta
40502 | 2022 | Wyoming, Uinta
40504 | 2022 | Wyoming, Washakie
72610 | 2030 | Wyoming, Washakie
72611 | 2030 | Wyoming, Washakie
72612 | 2030 | Wyoming, Washakie
72613 | 2030 | Wyoming, Washakie
49856 | 2026 | Wyoming, Washakie
40503 | 2022 | Wyoming, Washakie
49855 | 2026 | Wyoming, Washakie
49854 | 2026 | Wyoming, Washakie
40505 | 2022 | Wyoming, Washakie
12448 | 2010 | Wyoming, Washakie
12447 | 2010 | Wyoming, Washakie
12445 | 2010 | Wyoming, Washakie
31152 | 2018 | Wyoming, Washakie
31153 | 2018 | Wyoming, Washakie
21803 | 2014 | Wyoming, Washakie
21802 | 2014 | Wyoming, Washakie
31154 | 2018 | Wyoming, Washakie
21801 | 2014 | Wyoming, Washakie
12446 | 2010 | Wyoming, Washakie
12452 | 2010 | Wyoming, Weston
12449 | 2010 | Wyoming, Weston
31156 | 2018 | Wyoming, Weston
31155 | 2018 | Wyoming, Weston
12451 | 2010 | Wyoming, Weston
21806 | 2014 | Wyoming, Weston
21805 | 2014 | Wyoming, Weston
21804 | 2014 | Wyoming, Weston
12450 | 2010 | Wyoming, Weston
31157 | 2018 | Wyoming, Weston
49859 | 2026 | Wyoming, Weston
72616 | 2030 | Wyoming, Weston
40507 | 2022 | Wyoming, Weston
49858 | 2026 | Wyoming, Weston
49857 | 2026 | Wyoming, Weston
72615 | 2030 | Wyoming, Weston
40506 | 2022 | Wyoming, Weston
72617 | 2030 | Wyoming, Weston
72614 | 2030 | Wyoming, Weston
40508 | 2022 | Wyoming, Weston
(72617 rows)

```

Single quotes denote string constants. Double quotes denote things like tables and column names. If you want an upper-lower case or a table name or column name with blanks then you have to quote it with double quotes.

Lot's of stuff will fail if you put blanks in your table names. This is worse than putting blanks in your file names. Python will crash with blanks in file names.

We will be using more than one table in queries. To do this we need to tell SQL the table. That is done with a table-alias. In this case *t1*.

File: 06.sql

```
select t1.id, t1.year + 10 as "x", t1.state||', '||t1.county_name as "Location"
      from vote_by_county as t1
      order by 3 asc
;
```

There are lots of builtin functions that you can use and all sorts of arithmetic operators in PostgreSQL that can be applied to the projected columns. You can also write your own functions and pass data from the projected columns to the function and get back results.

Languages for functions include PG/SQL - the built in PostgreSQL language and others like JavaScript, Lua, C, C++, Go etc. I use PG/SQL and C for processing. JavaScript is 5 to 10x slower than PG/SQL. C is 10x faster than PG/SQL. These are rough numbers. I am a big fan of Go but I haven't used it for stored-procedure/functions yet in PostgreSQL.

Using a C function requires re-loading and re-starting the database - so it is hard.

Operators

<https://www.postgresql.org/docs/9.0/functions.html>

There are lots!

Operator	Description	Example	Result
+	addition	2 + 3	5
-	subtraction	2 - 3	-1
*	multiplication	2 * 3	6
/	division (integer division truncates the result)	4 / 2	2
%	modulo (remainder)	5 % 4	1
^	exponentiation	2.0 ^ 3.0	8
/	square root	/ 25.0	5
/	cube root	/ 27.0	3
!	factorial	5 !	120
!!	factorial (prefix operator)	!! 5	120
@	absolute value	@ -5.0	5
&	bitwise AND	91 & 15	11
	bitwise OR	32 3	35
#	bitwise XOR	17 # 5	20
~	bitwise NOT	~1	-2

Operator	Description	Example	Result
<<	bitwise shift left	1 << 4	16
>>	bitwise shift right	8 >> 2	2

and string operations

Function	Return Type	Description	Example	Result
<code>string string</code>	text	String concatenation	<code>'Post' 'greSQL'</code>	PostgreSQL
<code>string non-string</code> or <code>non-string string</code>	text	String concatenation with one non-string input	<code>'Value: ' 42</code>	Value: 42
<code>bit_length(string)</code>	int	Number of bits in string	<code>bit_length('jose')</code>	32
<code>char_length(string)</code> or <code>character_length(string)</code>	int	Number of characters in string	<code>char_length('jose')</code>	4
<code>lower(string)</code>	text	Convert string to lower case	<code>lower('TOM')</code>	tom
<code>octet_length(string)</code>	int	Number of bytes in string	<code>octet_length('jose')</code>	4
<code>overlay(string placing string from int [for int])</code>	text	Replace substring	<code>overlay('Txxxxas' placing 'hom' from 2 for 4)</code>	Thomas
<code>position(substring in string)</code>	int	Location of specified substring	<code>position('om' in 'Thomas')</code>	3
<code>substring(string [from int] [for int])</code>	text	Extract substring	<code>substring('Thomas' from 2 for 3)</code>	hom
<code>substring(string from pattern)</code>	text	Extract substring matching POSIX regular expression. See Section 9.7 for more information on pattern matching.	<code>substring('Thomas' from '...\$')</code>	mas
<code>substring(string from pattern for escape)</code>	text	Extract substring matching SQL regular expression. See Section 9.7 for more information on pattern matching.	<code>substring('Thomas' from '%#o_a#"' for '#')</code>	oma
<code>trim([leading trailing both] [characters] from string)</code>	text	Remove the longest string containing only the characters (a space by default) from the start/end/both ends of the string	<code>trim(both 'x' from 'xTomxx')</code>	Tom
<code>upper(string)</code>	text	Convert string to upper case	<code>upper('tom')</code>	TOM

Base Functions

Function	Return Type	Description	Example	Result

Function	Return Type	Description	Example	Result
<code>abs(x)</code>	(same as input)	absolute value	<code>abs(-17.4)</code>	17.4
<code>cbrt(dp)</code>	dp	cube root	<code>cbrt(27.0)</code>	3
<code>ceil(dp or numeric)</code>	(same as input)	smallest integer not less than argument	<code>ceil(-42.8)</code>	-42
<code>ceiling(dp or numeric)</code>	(same as input)	smallest integer not less than argument (alias for <code>ceil</code>)	<code>ceiling(-95.3)</code>	-95
<code>degrees(dp)</code>	dp	radians to degrees	<code>degrees(0.5)</code>	28.6478897565412
<code>div(y numeric, x numeric)</code>	numeric	integer quotient of y/x	<code>div(9,4)</code>	2
<code>exp(dp or numeric)</code>	(same as input)	exponential	<code>exp(1.0)</code>	2.71828182845905
<code>floor(dp or numeric)</code>	(same as input)	largest integer not greater than argument	<code>floor(-42.8)</code>	-43
<code>ln(dp or numeric)</code>	(same as input)	natural logarithm	<code>ln(2.0)</code>	0.693147180559945
<code>log(dp or numeric)</code>	(same as input)	base 10 logarithm	<code>log(100.0)</code>	2
<code>log(b numeric, x numeric)</code>	numeric	logarithm to base b	<code>log(2.0, 64.0)</code>	6.0000000000
<code>mod(y, x)</code>	(same as argument types)	remainder of y/x	<code>mod(9,4)</code>	1
<code>pi()</code>	dp	" π " constant	<code>pi()</code>	3.14159265358979
<code>power(a dp, b dp)</code>	dp	a raised to the power of b	<code>power(9.0, 3.0)</code>	729
<code>power(a numeric, b numeric)</code>	numeric	a raised to the power of b	<code>power(9.0, 3.0)</code>	729
<code>radians(dp)</code>	dp	degrees to radians	<code>radians(45.0)</code>	0.785398163397448
<code>round(dp or numeric)</code>	(same as input)	round to nearest integer	<code>round(42.4)</code>	42
<code>round(v numeric, s int)</code>	numeric	round to s decimal places	<code>round(42.4382, 2)</code>	42.44
<code>sign(dp or numeric)</code>	(same as input)	sign of the argument (-1, 0, +1)	<code>sign(-8.4)</code>	-1
<code>sqrt(dp or numeric)</code>	(same as input)	square root	<code>sqrt(2.0)</code>	1.4142135623731
<code>trunc(dp or numeric)</code>	(same as input)	truncate toward zero	<code>trunc(42.8)</code>	42
<code>trunc(v numeric, s int)</code>	numeric	truncate to s decimal places	<code>trunc(42.4382, 2)</code>	42.43

Function	Return Type	Description	Example	Result
<code>width_bucket(numeric, b1 numeric, b2 numeric, count int)</code>	int	return the bucket to which operand would be assigned in an equidepth histogram with count buckets, in the range b1 to b2	<code>width_bucket(5.35, 0.024, 10.06, 5)</code>	3
<code>width_bucket(op dp, b1 dp, b2 dp, count int)</code>	int	return the bucket to which operand would be assigned in an equidepth histogram with count buckets, in the range b1 to b2	<code>width_bucket(5.35, 0.024, 10.06, 5)</code>	3

How About the square root operator!

File: 07.sql

```
select |/25.0;
```

and a factorial operator!

```
select 5!;
```

and

```
select !! 5;
```

That is *FUN* ! Not 1 but 2 factorial operators.